

# Artificial Intelligence

## Assignment – 2

**Name: C. Deepthi Chowdary**

**Reg no.: 20BCE2445**

```
[64]: import pandas as pd import numpy as np from  
sklearn.preprocessing import LabelEncoder,  
StandardScaler from sklearn.model_selection import  
train_test_split
```

```
[65]: df = pd.read_csv('drug200.csv')
```

```
[66]: # Task 1 : Read the dataset and do data pre-processing
```

```
[67]: label_encoder = LabelEncoder() df['Sex'] =  
label_encoder.fit_transform(df['Sex']) df['BP'] =  
label_encoder.fit_transform(df['BP']) df['Cholesterol'] =  
label_encoder.fit_transform(df['Cholesterol']) df['Drug']  
= label_encoder.fit_transform(df['Drug'])  
print(df.head())
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	0	0	0	25.355	0
1	47	1	1	0	13.093	3
2	47	1	1	0	10.114	3
3	28	0	2	0	7.798	4

```
4    61    0    1            0    18.043    0
```

```
[68]: # Scale numerical variables
```

```
scaler = StandardScaler()
df[['Age', 'Na_to_K']] = scaler.fit_transform(df[['Age', 'Na_to_K']])
```

```
[69]: # Separate features and labels
```

```
x = df[['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K']]
y = df['Drug']
```

```
[70]: # Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
print(X_train.shape)
print(y_test.shape)
```

```
(160, 5)
```

```
(40,)
```

```
[71]: # Task 2 : Build the ANN model with (input layer, min 3 hidden layers & output
      ↪ layer)
```

```
[72]: import tensorflow as tf
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense
```

```
[73]: # Define the model architecture
```

```
model = Sequential()
model.add(Dense(64, activation='relu', input_shape=(5,)))
model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(5, activation='softmax'))
```

```
[74]: x = df.iloc[:,0:5]
      y = df.iloc[:,5:]
      print(x)
      print(y)
```

	Age	Sex	BP	Cholesterol	Na_to_K
0	-1.291591	0	0	0	1.286522
1	0.162699	1	1	0	-0.415145

```

2    0.162699    1    1    0 -0.828558
3   -0.988614    0    2    0 -1.149963
4    1.011034    0    1    0  0.2717941
2
3

```

**Error! Bookmark not defined.**

8  
8

```

..      ... .. ..      ...      ...
195  0.708057    0    1    0 -0.626917
196 -1.715759    1    1    0 -0.565995
197  0.465676    1    2    0 -0.859089
198 -1.291591    1    2    1 -0.286500
199 -0.261469    0    1    1 -0.657170

```

[200 rows x 5 columns]

```

Drug
0    0
4    0
..   ...
195   3
196   3
197   4
198   4
199   4

```

[200 rows x 1 columns]

[75]: *# Compile the model*

```

model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam', metrics=['accuracy'])

y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)
model.fit(X_train, y_train_encoded, epochs=20, batch_size=20,
        validation_data=(X_test, y_test_encoded))

```

Epoch 1/20

8/8 [=====] - 2s 38ms/step - loss: 1.4517 - accuracy:

0.5813 - val\_loss: 1.3748 - val\_accuracy: 0.4000

Epoch 2/20

8/8 [=====] - 0s 6ms/step - loss: 1.2047 - accuracy:

0.5375 - val\_loss: 1.1855 - val\_accuracy: 0.4250  
Epoch 3/20  
8/8 [=====] - 0s 9ms/step - loss: 1.0034 - accuracy:  
0.6187 - val\_loss: 1.0329 - val\_accuracy: 0.5750  
Epoch 4/20  
8/8 [=====] - 0s 9ms/step - loss: 0.8368 - accuracy:  
0.7188 - val\_loss: 0.8926 - val\_accuracy: 0.6250  
Epoch 5/20  
8/8 [=====] - 0s 6ms/step - loss: 0.7157 - accuracy:  
0.7188 - val\_loss: 0.8098 - val\_accuracy: 0.6250  
Epoch 6/20  
8/8 [=====] - 0s 8ms/step - loss: 0.6184 - accuracy:  
0.7500 - val\_loss: 0.7295 - val\_accuracy: 0.7250  
Epoch 7/20  
8/8 [=====] - 0s 6ms/step - loss: 0.5321 - accuracy:  
0.8125 - val\_loss: 0.6841 - val\_accuracy: 0.7500  
Epoch 8/20  
8/8 [=====] - 0s 10ms/step - loss: 0.4566 - accuracy:  
0.8687 - val\_loss: 0.6015 - val\_accuracy: 0.8500  
Epoch 9/20  
8/8 [=====] - 0s 7ms/step - loss: 0.3843 - accuracy:  
0.9062 - val\_loss: 0.5173 - val\_accuracy: 0.8750  
Epoch 10/20  
8/8 [=====] - 0s 8ms/step - loss: 0.3252 - accuracy:  
0.9125 - val\_loss: 0.4404 - val\_accuracy: 0.8750  
Epoch 11/20  
8/8 [=====] - 0s 7ms/step - loss: 0.2629 - accuracy:  
0.9125 - val\_loss: 0.3672 - val\_accuracy: 0.8750  
Epoch 12/20  
8/8 [=====] - 0s 7ms/step - loss: 0.2216 - accuracy:  
0.9312 - val\_loss: 0.3321 - val\_accuracy: 0.8750  
Epoch 13/20  
8/8 [=====] - 0s 7ms/step - loss: 0.1819 - accuracy:  
0.9438 - val\_loss: 0.2550 - val\_accuracy: 0.9000  
Epoch 14/20

```

8/8 [=====] - 0s 10ms/step - loss: 0.1560 -
accuracy:
0.9500 - val_loss: 0.2532 - val_accuracy: 0.9500
Epoch 15/20
8/8 [=====] - 0s 6ms/step - loss: 0.1443 -
accuracy:
0.9688 - val_loss: 0.1985 - val_accuracy: 0.9000
Epoch 16/20
8/8 [=====] - 0s 9ms/step - loss: 0.1254 -
accuracy:
0.9688 - val_loss: 0.1833 - val_accuracy: 0.9750
Epoch 17/20
8/8 [=====] - 0s 6ms/step - loss: 0.0970 -
accuracy:
0.9875 - val_loss: 0.1717 - val_accuracy: 1.0000
Epoch 18/20
8/8 [=====] - 0s 9ms/step - loss: 0.0868 -
accuracy:
0.9750 - val_loss: 0.1504 - val_accuracy: 0.9750
Epoch 19/20
8/8 [=====] - 0s 6ms/step - loss: 0.0766 -
accuracy:
1.0000 - val_loss: 0.1436 - val_accuracy: 0.9750
Epoch 20/20
8/8 [=====] - 0s 7ms/step - loss: 0.0678 -
accuracy:
0.9812 - val_loss: 0.1206 - val_accuracy: 0.9750

```

```
[75]: <keras.callbacks.History at 0x7fc722a7be20>
```

```
[76]: y_pred = model.predict(x_test)
      y_pred
```

```

WARNING:tensorflow:5 out of the last 7 calls to <function
Model.make_predict_function.<locals>.predict_function at
0x7fc722bf49d0> triggered tf.function retracing. Tracing is
expensive and the excessive number of tracings could be due to (1)
creating @tf.function repeatedly in a loop, (2) passing tensors with
different shapes, (3) passing Python objects instead of tensors. For
(1), please define your @tf.function outside of the loop. For (2),
@tf.function has reduce_retracing=True option that can avoid
unnecessary retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling_retracing and
https://www.tensorflow.org/api_docs/python/tf/function for more
details.
2/2 [=====] - 0s 9ms/step

```

```
[76]: array([[4.13405127e-04, 1.27605614e-04, 2.03855492e-07,
7.50870770e-03, 9.91949975e-01],
[9.94201958e-01, 5.14725503e-03, 2.99533876e-05, 4.84759919e-
04, 1.36094895e-04],
[2.79626124e-06, 1.99977421e-06, 5.16646413e-11, 6.72629918e-
04, 9.99322474e-01],
[2.83280946e-03, 3.48852053e-02, 8.92015360e-03, 7.59812355e-
01, 1.93549350e-01],
[9.99999940e-01, 3.28292191e-19, 1.42062910e-17, 8.46457494e-
17, 5.58904698e-17],
[9.99691248e-01, 2.56415988e-05, 2.51631485e-04, 2.94335568e-
05, 2.17517095e-06],
[9.99999940e-01, 3.61117553e-10, 4.05409484e-10, 1.11134280e-
09, 9.09846420e-10],
[7.46123632e-03, 1.53253040e-05, 2.05253734e-08, 1.85971186e-
02, 9.73926246e-01],
[4.89533022e-02, 8.14404786e-01, 6.96765035e-02, 5.54476641e-
02, 1.15178749e-02],
[3.14717290e-05, 3.12856696e-06, 1.03769771e-07, 3.07339523e-
03, 9.96891856e-01],
[8.33706290e-04, 9.44750011e-01, 4.69562830e-03, 4.91494723e-
02, 5.71190671e-04],
[5.63477771e-03, 1.65499118e-03, 4.97897986e-07, 2.14239918e-
02, 9.71285701e-01],
[9.99937952e-01, 3.12065737e-07, 1.05881973e-07, 2.78759489e-
05, 3.36685516e-05],
[3.92728811e-03, 9.50904250e-01, 2.91301263e-03, 4.14308533e-
02, 8.24655988e-04],
[2.11916384e-04, 1.94486752e-02, 9.77127016e-01, 3.20940185e-
03, 2.85138822e-06],
[9.99988854e-01, 2.64876510e-10, 1.12958193e-11, 9.42942393e-
07, 1.01327441e-05],
[1.60759955e-03, 1.64753329e-02, 9.78582621e-01, 3.29385232e-
03, 4.04419807e-05],
[1.57631177e-06, 4.22669018e-07, 7.01798897e-10, 9.63229686e-
04, 9.99034703e-01],
[3.98420263e-04, 1.10615864e-01, 1.25297796e-04, 5.62819958e-
01, 3.26040477e-01],
[9.99999940e-01, 2.10215739e-14, 7.02131292e-14, 1.55016607e-
11, 5.87058735e-11],
[8.40014219e-03, 1.10281460e-01, 8.65873754e-01, 1.37768965e-
02, 1.66778930e-03],
[5.21895364e-02, 9.92505578e-04, 2.03632610e-03, 1.45251110e-
01, 7.99530506e-01],
[8.76396836e-04, 2.67904103e-02, 9.21104662e-03, 4.60485995e-
01,
```

```

5.02636135e-01]],
[9.99999940e-01, 6.66354848e-15, 7.17282204e-14, 6.50112885e-
13, 1.00215138e-12],
[9.99999940e-01, 5.00953337e-16, 5.93842814e-15, 3.52168192e-
13, 5.14562525e-12],
[9.99999940e-01, 2.64196543e-15, 2.55897327e-14, 2.75578768e-
13, 3.84631481e-13],
[1.00730290e-03, 5.72257526e-02, 1.34035340e-03, 6.65092647e-
01, 2.75333911e-01],
[2.08249821e-05, 4.83725955e-07, 1.95186818e-11, 9.81732621e-
04, 9.98996973e-01],
[9.99999940e-01, 3.62774255e-11, 6.37677827e-11, 1.92503111e-
10, 1.42245091e-10],
[1.29936814e-01, 4.21307086e-05, 3.51125891e-06, 8.77872203e-
03, 8.61238778e-01],
[9.99990046e-01, 5.69632475e-09, 3.74583742e-09, 6.74399985e-
07, 9.22276013e-06],
[1.28411793e-05, 1.30465448e-01, 7.51612561e-06, 8.05001497e-
01, 6.45127445e-02],
[1.78256020e-01, 1.00485990e-02, 5.48207936e-05, 3.79701257e-
01, 4.31939214e-01],
[9.9999583e-01, 2.13776746e-10, 6.09901921e-11, 3.22761871e-
08, 3.09697043e-07],
[1.15087496e-04, 8.31787109e-01, 1.56512201e-01, 1.14013907e-
02, 1.84151490e-04],
[9.99999940e-01, 6.83683931e-14, 4.79056085e-13, 1.24546218e-
12, 5.17932702e-13],
[1.88411415e-01, 1.24890450e-03, 5.95483556e-03, 1.63057938e-
01, 6.41326845e-01],
[2.12751655e-03, 9.30602849e-01, 2.18930449e-02, 4.26748469e-
02, 2.70170020e-03],
[9.99997914e-01, 6.68790108e-07, 2.85858519e-08, 8.10713004e-
07, 4.23714482e-07],
[4.69133374e-04, 9.55850482e-01, 1.62037276e-02, 2.63245087e-
02,
1.15206011e-03]], dtype=float32)

```

```

[77]: comp = pd.DataFrame(y_test_encoded) # Creating a
dataframe comp.columns = ['Actual Value'] # Changing
the column name comp

```

```

[77]: Actual Value
0      4
1      0

```

2	4
3	3
4	0
5	0
6	0
7	4
8	1
9	4
10	1
11	4
12	0
13	1
14	2
15	0
16	2
17	4
18	3
19	0
20	2
21	4
22	4
23	0
24	0
25	0
26	3
27	4
28	0
29	4
30	0
31	3
32	3
33	0
34	1
35	0
36	4
37	1
38	0
39	1

```
[78]: # Print the model summary
```

```
model.summary()
```

Model: "sequential\_1"

---

Layer (type)

Output Shape

Param #



```
=====
dense_5 (Dense)          (None, 64)          384
dense_6 (Dense)          (None, 128)         8320
dense_7 (Dense)          (None, 64)          8256
dense_8 (Dense)          (None, 32)          2080
dense_9 (Dense)          (None, 5)           165
=====
Total params: 19,205
Trainable params: 19,205
Non-trainable params: 0
```

```
[79]: # Task 3 : Test the model with random data
```

```
[80]: # Generate random data for testing
```

```
random_data = np.random.rand(1, 5)
random_data
```

```
[80]: array([[0.87039758, 0.52583504, 0.74177248, 0.71396893,
0.03728909]])
```

```
[81]: # Make predictions
```

```
predictions = model.predict(random_data)
predictions
```

WARNING:tensorflow:6 out of the last 9 calls to <function Model.make\_predict\_function.<locals>.predict\_function at 0x7fc722bf49d0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.

```
1/1 [=====] - 0s 77ms/step
```

```
[81]: array([[9.9052775e-01, 3.0603227e-05, 6.6905326e-05, 1.3001083e-03,
8.0746198e-03]], dtype=float32)
```

```
[82]: # Get the predicted drug class  
  
predicted_class = np.argmax(predictions)
```

```
[84]: # Print the predicted class  
  
print("Predicted Drug Class :", predicted_class)
```

Predicted Drug Class : 0